

CS 184A/284A: Artificial Intelligence in Biology and Medicine

Homework 2Due date: **Sunday, Oct 26, 2025**

Instructor: Xiaohui Xie

The submission for this homework, as for the first one, should be **one stand-alone PDF file** containing all of the relevant code, figures, and any text explaining your results. When coding the answers, try to write functions to encapsulate and reuse the code, instead of copy pasting the same code multiple times. This will not only reduce your programming efforts, but also make it easier to debug, and for us to grade your work.

Write neatly (or type) and show all your work!

Points: This homework adds up to a total of **100 points**, as follows:

Problem 1: Linear Regression	60 points
Problem 2: Cross-Validation	35 points
Statement of Collaboration	5 points

Problem 1: Linear Regression, (60 points)

For this problem we will explore linear regression, the creation of additional features, and cross-validation.

1. Load the “`data/curve80.txt`” data set, and split it into 75% / 25% training/test. The first column `data[:,0]` is the scalar feature (x) values; the second column `data[:,1]` is the target value y for each example. For consistency in our results, **don’t** reorder (shuffle) the data (they’re already in a random order), and use the first 75% of the data for training and the rest for testing:

```
1 import numpy as np
2 from sklearn.model_selection import train_test_split
3
4 data = np.genfromtxt("data/curve80.txt", delimiter=None)
5 X = data[:, 0].reshape(-1, 1)
6 Y = data[:, 1]
7 Xtr, Xte, Ytr, Yte = train_test_split(
8     X, Y, test_size=0.25, shuffle=False
9 )
```

Print the shapes of these four objects. **(5 points)**

2. Use `sklearn.linear_model.LinearRegression` to create a linear regression predictor of y given x . You can plot the resulting function by simply evaluating the model at a large number of x values, `xs`:

```
1 from sklearn.linear_model import LinearRegression
2 lr = LinearRegression()
3 lr.fit(Xtr, Ytr)
4 xs = np.linspace(0,10,200).reshape(-1,1)
5 ys = lr.predict(xs)
```

- (a) Plot the training data points along with your prediction function in a single plot. Plot the ground truth testing data points in a different color and marker shape. **(10 points)**
- (b) Print the linear regression coefficients (`lr.coef_, lr.intercept_`) and verify that they match your plot. **(5 points)**
- (c) What is the mean squared error of the predictions on the training and test data? **(10 points)**

3. Try fitting $y = f(x)$ using a polynomial function $f(x)$ of increasing order. Do this by the trick of adding additional polynomial features before constructing and training the linear regression object. You can do this easily yourself; you can add a quadratic feature of `Xtr` with

```
1 Xtr2 = np.zeros( (Xtr.shape[0],2) ) # create Mx2 array to store features
2 Xtr2[:,0] = Xtr[:,0] # place original "x" feature as X1
3 Xtr2[:,1] = Xtr[:,0]**2 # place "x^2" feature as X2
4 # Now, Xtr2 has two features about each data point: "x" and "x^2"
```

(You can also add the all-ones constant feature in a similar way, but this is currently done automatically within the learner's train function.) `sklearn.preprocessing.PolynomialFeatures` can be used to create such features. Note, though, that the resulting features may include extremely large values – if $x \approx 10$, then e.g., x^{10} is extremely large. For this reason (as is often the case with features on very different scales) it's a good idea to rescale the features; again, you can do this manually or use `sklearn.preprocessing.StandardScaler`:

```
1 from sklearn.preprocessing import PolynomialFeatures, StandardScaler
2
3 poly = PolynomialFeatures(degree=degree, include_bias=False)
4 scaler = StandardScaler()
5 Xtr_poly = poly.fit_transform(Xtr) # create polynomial features
6 Xtr_poly = scaler.fit_transform(Xtr_poly) # scale features
7 Xte_poly = poly.transform(Xte) # create polynomial features
8 Xte_poly = scaler.transform(Xte_poly) # scale features
9
10 lr = LinearRegression()
11 lr.fit(Xtr_poly, Ytr) # fit model
```

This snippet also shows a useful feature transformation framework – often we wish to apply some transformation to the features; in many cases the desired transformation depends on the data (such as rescaling the data to unit variance). Ideally, we should then be able to apply this same transform to new test data when it arrives, so that it will be treated in exactly the same way as the training data.

Train models of degree $d = 1, 3, 5, 7, 10, 18$ and:

- Plot their learned prediction functions $f(x)$ (15 points)
- Plot their training and test errors on a log scale (`semilogy`) as a function of the degree. (10 points)
- What polynomial degree do you recommend? Do you see any trends regarding train and test performance as a function of the degree? (5 points)

For (a), remember that your learner has now been trained on the polynomially expanded features, and so is expecting `degree` features (columns) to be input. So, don't forget to also expand and scale the features of `xs` as well. You can do this manually as in the code snippet above, or you can think of this as a "feature transform" function phi, eg.,

```
1 def phi(X, poly, scaler):
2     return scaler.transform(poly.transform(X))
3
4 # Now, phi will do the required feature expansion and rescaling:
5 YhatTrain = lr.predict(phi(Xtr, poly, scaler)) # predict on training data
6 YhatTest = lr.predict(phi(Xte, poly, scaler)) # predict on test data
```

Also, you may want to save the original axes of your plot and re-apply them to each subsequent plot for consistency. (Otherwise, high-degree polynomials may look "flat" due to some extremely large values.) You can do this as shown in Discussion 1 notebook by, for example:

```
1 # Creating subplots with just one subplot so basically a single figure.
2 fig, ax = plt.subplots(1, 1, figsize=(10, 8))
3 ax.plot(...) # Plot for each polynomial degree
4 ax.plot(...) # like so
5 ax.set_ylim(..., ...) # Set the minimum and maximum limits
6 plt.show()
```

Problem 2: Cross-validation (35 points)

In the previous problem, you decided what degree of polynomial fit to use based on performance on some test data¹. Let's now imagine that you did not have access to the target values of the test data you held out in the previous problem, and wanted to decide on the best polynomial degree.

Of course, we could simply repeat the exercise, further splitting `Xtr` into a training and validation split, and then assessing performance on the validation data to decide on a degree. But when training is reasonably fast, it can be more effective to use cross-validation to estimate the optimal degree. Cross-validation works by creating many such training/validation splits, called folds, and using all of these splits to assess the “out-of-sample” (validation) performance by averaging them. You can do a 5-fold validation test, for example, by:

```
1 from sklearn.model_selection import KFold
2
3 n_splits = 5
4 kf = KFold(n_splits=n_splits, shuffle=True, random_state=42)
5 val_mse_list = []
6
7 for train_index, val_index in kf.split(Xtr):
8     Xti, Xvi = Xtr[train_index], Xtr[val_index]
9     Yti, Yvi = Ytr[train_index], Ytr[val_index]
10    # TODO: preprocess Xti, Xvi
11    # TODO: train model on Xti, Yti
12    # TODO: predict on Xvi
13    val_mse = # TODO: compute validation MSE
14    val_mse_list.append(val_mse)
15 # the overall estimated validation error is the average of the error on each fold
16 print(np.mean(val_mse_list))
```

Using this technique on your training data `Xtr` from the previous problem, find the 5-fold cross-validation MSE of linear regression at the same degrees as before, $d = 1, 3, 5, 7, 10, 18$ (or more densely, if you prefer). Again, a function that has degree and number of folds as arguments, and returns cross-validation error, will be useful.

1. Plot the **five-fold** cross-validation error and test error (with **semilogy**, as before) as a function of degree. (10 points)
2. How do the MSE estimates from five-fold cross-validation compare to the MSEs evaluated on the actual test data (Problem 1)? (5 points)
3. Which polynomial degree do you recommend based on five-fold cross-validation error? (5 points)
4. For the degree that you picked in step 3, plot the cross-validation error as the number of folds is varied ($n_{\text{Folds}} = 2, 3, 4, 5, 6, 10, 12, 15$), again with **semilogy**. What pattern do you observe, and how do you explain it? (15 points)

Statement of Collaboration (5 points)

It is **mandatory** to include a **Statement of Collaboration** in each submission, with respect to the guidelines below. Include the names of everyone involved in the discussions (especially in-person ones), and what was discussed.

All students are required to follow the academic honesty guidelines posted on the course website. For programming assignments, in particular, I encourage the students to organize (perhaps using Campuswire) to discuss the task descriptions, requirements, bugs in my code, and the relevant technical content **before** they start working on it. However, you should not discuss the specific solutions, and, as a guiding principle, you are not allowed to take anything written or drawn away from these discussions (i.e. no photographs of the blackboard, written notes, referring to Campuswire, etc.). Especially **after** you have started working on the assignment, try to restrict the discussion to Campuswire as much as possible, so that there is no doubt as to the extent of your collaboration.

¹Technically, since you knew the answers to these data's targets and could use them to evaluate performance at different degrees, I would probably call them validation data instead.